

# Mudcard

- **loc and iloc :**)
  - we will go through this again today
- **should we just use pandas in this class or would using both/polars benefit us more in and out of the classroom?**
  - I'd say learn both, that's the most beneficial for you out of classroom.
  - As far as problem sets and quizzes are concerned, it's fine to use one of the packages.
  - I'll mostly use pandas going forward.
- **What are some ways to get more familiar with working with python?**
  - easy leetcode problems are a good start to learn standard python (no packages)
  - you will practice how to work with popular data science packages well enough during class and the problem sets I think
- **need to review basic python functions, like skiprows**
  - skiprows is an argument of the pd.read\_csv function :)

## Lecture 3

### Working with data (step 0) continued

### Learning objectives

By the end of the lecture, you will be able to

- filter columns
- merge and append data frames
- modify a dataframe

## Pandas and Polars

- data are often distributed over multiple files/databases (e.g., csv and excel files, sql databases)
- each file/database is read into a pandas dataframe - last lecture
- you often need to filter dataframes by selecting specific rows - last lecture
- you often need to filter dataframes by selecting specific columns - today
- multiple dataframes need to be merged and appended - today
- dataframes sometimes need to be modified - today

## Some notes and advice

- **ALWAYS READ THE HELP OF THE METHODS/FUNCTIONS YOU USE!**
- stackoverflow is your friend, use it! <https://stackoverflow.com/>
- you can also use generative AI to help you fix bugs
  - Gemini is supported by Brown OIT (see here <https://go.brown.edu/gemini>)
- [here](#) is an excellent review of the syntax differences between pandas and polars

## Learning objectives

By the end of the lecture, you will be able to

- **filter columns**
- merge and append data frames
- modify a dataframe

```
In [1]: import pandas as pd
df_pd = pd.read_csv('../data/adult_data.csv')

columns = df_pd.columns
print(columns)

# select columns by column name
#print(df_pd[['age', 'hours-per-week']])
#print(columns[[1,5,7]])
#print(df_pd[columns[[1,5,7]]])

# select columns by index using iloc
#print(df_pd.iloc[:,3])

# select columns by index - not standard python indexing
#print(df_pd.iloc[:, [3,5,6]])

# select columns by index - standard python indexing
#print(df_pd.iloc[:, ::2])
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
      'marital-status', 'occupation', 'relationship', 'race', 'sex',
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
      'gross-income'],
      dtype='object')
```

```
In [2]: import polars as pl
df_pl = pl.read_csv('../data/adult_data.csv')

columns = df_pl.columns
print(columns)

# select columns by column name
#print(df_pl['age', 'hours-per-week'])
```

```
#print(df_pl.select(['age','hours-per-week'])) # use .select if you know the
#print(columns[1:4]) # indices must be integers or slices
#print(df_pl[columns[1:4]])

# select columns by index, polars has no .iloc
#print(df_pl[:,3])

# select columns by index - not standard python indexing but it works
#print(df_pl[:,[3,5,6]])

# select columns by index - standard python indexing
#print(df_pl[:,,:2])
```

```
['age', 'workclass', 'fnlwt', 'education', 'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'gross-income']
```

## Learning objectives

By the end of the lecture, you will be able to

- filter columns
- **merge and append data frames**
- modify a dataframe

## How to merge dataframes in Pandas?

Merge - info on data points are distributed in multiple files

In [3]: *# We have two datasets from two hospitals*

```
hospital1 = {'ID':['ID1','ID2','ID3','ID4','ID5','ID6','ID7'],'col1':[5,8,2,
df1 = pd.DataFrame(data=hospital1)
print(df1)

hospital2 = {'ID':['ID2','ID5','ID6','ID10','ID11'],'col3':[12,76,34,98,65],
df2 = pd.DataFrame(data=hospital2)
print(df2)
```

	ID	col1	col2
0	ID1	5	y
1	ID2	8	j
2	ID3	2	w
3	ID4	6	b
4	ID5	0	a
5	ID6	2	b
6	ID7	5	t

  

	ID	col3	col2
0	ID2	12	q
1	ID5	76	u
2	ID6	34	e
3	ID10	98	l
4	ID11	65	p

```
In [4]: # we are interested in only patients from hospital1
df_left = df1.merge(df2,how='left',on='ID') # IDs from the left dataframe (c
#print(df_left)

# we are interested in only patients from hospital2
#df_right = df1.merge(df2,how='right',on='ID') # IDs from the right datafram
#df_right = df2.merge(df1,how='left',on='ID')
#print(df_right)

# we are interested in patients who were in both hospitals
#df_inner = df1.merge(df2,how='inner',on='ID') # merging on IDs present in b
#print(df_inner)

# we are interested in all patients who visited at least one of the hospital
#df_outer = df1.merge(df2,how='outer',on='ID') # merging on IDs present in
#print(df_outer)
```

## How to append dataframes in pandas?

Append - new data comes in over a period of time. E.g., one file per month/quarter/fiscal year etc.

You want to combine these files into one data frame.

```
In [5]: df_append = pd.concat([df1,df2]) # note that rows with ID2, ID5, and ID6 ar
print(df_append)

#df_append = pd.concat([df1,df2],ignore_index=True) # note that rows with ID
#print(df_append)

# d3 = {'ID':['ID23','ID94','ID56','ID17'],'col1':['rt','h','st','ne'],'col2
# df3 = pd.DataFrame(data=d3)
# print(df3)

# df_append = pd.concat([df1,df2,df3],ignore_index=True) # multiple datafram
# print(df_append)
```

	ID	col1	col2	col3
0	ID1	5.0	y	NaN
1	ID2	8.0	j	NaN
2	ID3	2.0	w	NaN
3	ID4	6.0	b	NaN
4	ID5	0.0	a	NaN
5	ID6	2.0	b	NaN
6	ID7	5.0	t	NaN
0	ID2	NaN	q	12.0
1	ID5	NaN	u	76.0
2	ID6	NaN	e	34.0
3	ID10	NaN	l	98.0
4	ID11	NaN	p	65.0

## How to merge/join dataframes in Polars?

```
In [6]: hospital1 = {'ID': ['ID1', 'ID2', 'ID3', 'ID4', 'ID5', 'ID6', 'ID7'], 'col1': [5, 8, 2,
df1 = pl.DataFrame(data=hospital1)
#print(df1)

hospital2 = {'ID': ['ID2', 'ID5', 'ID6', 'ID10', 'ID11'], 'col3': [12, 76, 34, 98, 65],
df2 = pl.DataFrame(data=hospital2)
#print(df2)

# left join
#df_left = df1.join(df2, how='left', on='ID') # IDs from the left dataframe (c
#print(df_left)

# right join
#df_right = df1.join(df2, how='right', on='ID') # IDs from the right dataframe
#df_right = df2.join(df1, how='left', on='ID')
#print(df_right)

# inner join
#df_inner = df1.join(df2, how='inner', on='ID') # merging on IDs present in bo
#print(df_inner)

# outer join is called a full join
#df_outer = df1.join(df2, how='full', on='ID') # merging on IDs present in ar
#print(df_outer)
```

## Quiz

```
In [7]: raw_data_1 = {
    'subject_id': ['1', '2', '3', '4', '5'],
    'first_name': ['Alex', 'Amy', 'Allen', 'Alice', 'Ayoung'],
    'last_name': ['Anderson', 'Ackerman', 'Ali', 'Aoni', 'Atiches']}

raw_data_2 = {
    'subject_id': ['6', '7', '8', '9', '10'],
    'first_name': ['Billy', 'Brian', 'Bran', 'Bryce', 'Betty'],
    'last_name': ['Bonder', 'Black', 'Balwner', 'Brice', 'Btisan']}

raw_data_3 = {
    'subject_id': ['1', '2', '3', '4', '5', '7', '8', '9', '10', '11'],
    'test_id': [51, 15, 15, 61, 16, 14, 15, 1, 61, 16]}

# Create three data frames from raw_data_1, 2, and 3.
# Append the first two data frames and assign it to df_append.
# Merge the third data frame with df_append such that only subject_ids from
# Assign the new data frame to df_merge.
# How many rows and columns do we have in df_merge?
```

**Always check that the resulting dataframe is what you wanted to end up with!**

- small toy datasets are ideal to test your code.

If you need to do a more complicated dataframe operations, check out `pd.concat()` and `pl.concat()`!

## Learning objectives

By the end of the lecture, you will be able to

- filter columns
- merge and append data frames
- **modify a dataframe**

**Do not EVER overwrite the original data files!**

**Always save the modified dataframe to a new file!**

**Why do we need to modify the dataset?**

- feature engineering
  - generating new features (adding new informative columns) can improve the performance of the ML model
- fix dataset issues
  - there might be data quality issues that need to be manually fixed
  - typos, missing values, etc

```
In [8]: import pandas as pd
df_pd = pd.read_csv('../data/adult_data.csv')
print(df_pd.head())

# let's generate a new feature
# is immigrant? False (0) if the person's home country is the USA, True (1)
# such a feature has only two categories but it might be pretty informative
df_pd['is immigrant'] = df_pd['native-country'] != ' United-States'

#print(df_pd.head())
```

	age	workclass	fnlwt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

  

	marital-status	occupation	relationship	race	sex
0	Never-married	Adm-clerical	Not-in-family	White	Male
1	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

  

	capital-gain	capital-loss	hours-per-week	native-country	gross-income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

## Mud card

In [ ]: